

# Learning Combinatorial Optimization Algorithms over Graphs

Hanjun Dai, Elias Khalil, Yuyu Zhang, Bistra Dilkina, Le Song

NeurIPS'17

# Approach

**Input:** Graph  $G = (V, E)$ , weights  $w(u, v)$  for  $(u, v) \in E$

**Algorithm design pattern:** Greedy

Feasible solution constructed by successively adding nodes to solution

**State representation:** Graph embedding


**Algorithm training:** Fitted Q-learning

# Outline

- 1. Greedy algorithms**
2. Graph representation
3. RL formulation
4. Q-learning
5. Experiments

# Minimum vertex cover

Find smallest vertex subset such that each edge is covered




# Minimum vertex cover

Find smallest vertex subset such that each edge is covered

## **2-approximation:**

Greedily add vertices of edge with **maximum degree sum**




# Minimum vertex cover

Find smallest vertex subset such that each edge is covered

## 2-approximation:

Greedily add vertices of edge with maximum degree sum

**Scoring function** that guides greedy algorithm



# Maximum cut


Find partition  $(S, V \setminus S)$  of nodes that maximizes

$$\sum_{(u,v) \in C} w(u,v)$$

where  $C = \{(u,v) \in E : u \in S, v \notin S\}$

If  $w(u,v) = 1$  for all  $(u,v) \in E$ :

$$\sum_{(u,v) \in C} w(u,v) = 5$$



# Maximum cut


Find partition  $(S, V \setminus S)$  of nodes that maximizes

$$\sum_{(u,v) \in C} w(u,v)$$

where  $C = \{(u, v) \in E : u \in S, v \notin S\}$

**Greedy:** move node from one side of cut to the other

Move node that results in the largest improvement in cut weight





# Maximum cut

Find partition  $(S, V \setminus S)$  of nodes that maximizes


$$\sum_{(u,v) \in C} w(u,v)$$

where  $C = \{(u,v) \in E : u \in S, v \notin S\}$

**Greedy:** move node from one side of cut to the other

Move node that results in the largest improvement in cut weight

**Scoring function** that guides greedy algorithm




# General greedy algorithm formulation

1. **Partial solution** is an ordered list  $S = (v_1, v_2, \dots, v_{|S|})$ ,  $v_i \in V$
2. **Helper function**  $h(S)$  maps  $S$  to combinatorial structure, eg:
  - **Maxcut:**  $h(S)$  returns cut  $C = \{(u, v) \in E : u \in S, v \notin S\}$
  - **TSP:**  $h(S)$  maintains a partial tour according to order of nodes in  $S$
  - **Min vertex cover:**  $h(S)$  does nothing
3. **Quality** of  $S$  evaluated by function  $c(h(S), G)$ , e.g.:
  - **Maxcut:**  $c(h(S), G) = \sum_{(u,v) \in C=h(S)} w(u, v)$
  - **TSP:**  $c(h(S), G) = -\sum_{i=1}^{|S|-1} w(S[i], S[i+1]) - w(S[|S|], S(1))$
  - **Min vertex cover:**  $c(h(S), G) = -|S|$

# General greedy algorithm formulation

4. Add node that maximizes an evaluation function  $Q(h(S), v)$ :  
 $S \leftarrow (S, v^*)$  where  $v^* = \operatorname{argmax}_{v \notin S} Q(h(S), v)$
5. Terminate based on termination criterion  $t(h(S))$

**This paper:** Use RL to learn evaluation function  $\hat{Q}(h(S), v; \Theta)$



# Outline

1. Greedy algorithms
- 2. Graph representation**
3. RL formulation
4. Q-learning
5. Experiments

# Representation: graph embedding

- $x_v = \begin{cases} 1 & \text{if } v \in S \\ 0 & \text{else} \end{cases}$

- Compute embedding over  $T$  iterations ( $\mu_v^{(0)} = \mathbf{0}$ ):

$$\mu_v^{(t+1)} \leftarrow \text{relu} \left( \underbrace{\theta_1 x_v}_{\text{Trainable parameters}} + \underbrace{\theta_2 \sum_{u \in N(v)} \mu_u^{(t)}}_{\text{Trainable parameters}} + \underbrace{\theta_3 \sum_{u \in N(v)} \text{relu}(\theta_4 w(v, u))}_{\text{Trainable parameters}} \right)$$

Trainable parameters

# Representation: graph embedding

- $x_v = \begin{cases} 1 & \text{if } v \in S \\ 0 & \text{else} \end{cases}$
- Compute embedding over  $T$  iterations ( $\mu_v^{(0)} = \mathbf{0}$ ):  
$$\mu_v^{(t+1)} \leftarrow \text{relu} \left( \theta_1 x_v + \theta_2 \sum_{u \in N(v)} \mu_u^{(t)} + \theta_3 \sum_{u \in N(v)} \text{relu}(\theta_4 w(v, u)) \right)$$

(Usually  $T = 4$ )
- $\hat{Q}(h(S), v; \Theta) = \theta_5^\top \text{relu} \left( \underbrace{\left[ \theta_6 \sum_{u \in V} \mu_u^{(T)}, \theta_7 \mu_u^{(T)} \right]}_{\text{Concatenation}} \right)$

# Representation: graph embedding

- $x_v = \begin{cases} 1 & \text{if } v \in S \\ 0 & \text{else} \end{cases}$
- Compute embedding over  $T$  iterations ( $\mu_v^{(0)} = \mathbf{0}$ ):  
$$\mu_v^{(t+1)} \leftarrow \text{relu} \left( \theta_1 x_v + \theta_2 \sum_{u \in N(v)} \mu_u^{(t)} + \theta_3 \sum_{u \in N(v)} \text{relu}(\theta_4 w(v, u)) \right)$$

(Usually  $T = 4$ )
- $\hat{Q}(h(S), v; \Theta) = \theta_5^\top \text{relu} \left( \underbrace{\left[ \theta_6 \sum_{u \in V} \mu_u^{(T)} \right]}_{\text{Surrogate for } h(S)}, \underbrace{\theta_7 \mu_v^{(T)}}_{\text{Surrogate for } v} \right)$

# Outline

1. Greedy algorithms
2. Graph representation
- 3. RL formulation**
4. Q-learning
5. Experiments



# Reinforcement learning formulation

**State:**  $\sum_{u \in V} \mu_u^{(T)}$

**Action:** Choose vertex  $v \in V \setminus S$  to add to solution

**Transition** (deterministic): For chosen  $v \in V \setminus S$ , set  $x_v = 1$

# Reinforcement learning formulation

**Reward:**  $r(S, v)$  is objective change when move to  $S' = (S, v)$   
$$r(S, v) = c(h(S'), G) - c(h(S), G)$$

$c(h(\emptyset), G) = 0$ , so cumulative reward of **terminal state**  $\hat{S}$  is

$$\sum_{i=1}^{|\hat{S}|} r(S_i, v_i) = c(h(\hat{S}), G)$$

**Policy** (deterministic):  $\pi(v|S) = \begin{cases} 1 & \text{if } v = \operatorname{argmax}_{v' \notin S} \hat{Q}(h(S), v'; \Theta) \\ 0 & \text{else} \end{cases}$

# Outline

1. Greedy algorithms
2. Graph representation
3. RL formulation
- 4. Q-learning**
5. Experiments

# Q-learning

Recall standard (1-step) Q-learning:

$$\min_{\Theta} \left( y - \hat{Q}(h(S_t), v_t; \Theta) \right)^2$$

where  $y = r(S_t, v_t) + \gamma \max_{v'} \hat{Q}(h(S_{t+1}), v'; \Theta)$

## Challenge:

- Final objective value only revealed after many steps
- 1-step update may be too myopic

Instead, use **n-step** Q-learning [Watkins, '89]

# $n$ -step Q-learning

$$\min_{\Theta} \left( y - \hat{Q}(h(S_t), v_t; \Theta) \right)^2$$

where  $y = \sum_{i=0}^{n-1} \gamma^i r(S_{t+1}, v_{t+i}) + \gamma^n \max_{v'} \hat{Q}(h(S_{t+n}), v'; \Theta)$

# Q-learning for the greedy algorithm

initialize set  $M = \emptyset$

for episode  $e = 1, \dots, L$ :

    sample graph  $G$  from underlying distribution  $D$

    initialize state to empty  $S_1 = ()$

# Q-learning for the greedy algorithm

for episode  $e = 1, \dots, L$ :

for step  $t = 1, \dots, T$ :

$$v_t = \begin{cases} \text{random node } v \notin S_t & \text{with probability } \epsilon \\ \operatorname{argmax}_{v \notin S_t} \hat{Q}(h(S_t), v; \Theta) & \text{otherwise} \end{cases}$$

add  $v_t$  to partial solution  $S_{t+1} = (S_t, v_t)$

if  $t \geq n$ :

**experience  
replay** | add tuple  $(S_{t-n}, v_{t-n}, \sum_{i=1}^n R(S_{t-i}, v_{t-i}), S_t)$  to  $M$   
sample batch  $B \sim M$   
update  $\Theta$  using SGD over  $B$

# Outline

1. Greedy algorithms
2. Graph representation
3. RL formulation
4. Q-learning
- 5. Experiments**



# Approximation ratio

Results measured in terms of approximation ratio

$$\frac{\text{Algorithm's solution}}{\text{OPT}}$$

# Min vertex cover


Barabasi-Albert  
random graphs

Paper's approach

Another DL approach  
[Bello et al., arXiv'16]

2-approximation  
algorithm

Greedy algorithm  
from first few slides



# Max cut


Barabasi-Albert  
random graphs

Paper's approach

Another DL approach  
[Bello et al., arXiv'16]

Goemans-Williamson  
algorithm

Greedy algorithm  
from first few slides




# TSP

Uniform random points on 2-D grid

## Paper's approach


- Initial subtour: 2 cities that are farthest apart
- Repeat the following:
  - Choose city that's *farthest* from any city in the subtour
  - Insert in position where it causes the smallest distance increase

[Rosenkrantz et al., SIAM JoC'77]




# Runtime comparisons

**CPLEX-1st:** 1<sup>st</sup> feasible solution found by CPLEX



# Min vertex cover visualization



Nodes seem to be selected to balance between:

- Degree
- Connectivity of the remaining graph

# Overview

Learn greedy heuristics for hard combinatorial problem

Approach based on graph representation + RL

Suggest approach could be used for **algorithm discovery**

“New and interesting” greedy strategies

“which **intuitively make sense** but have **not been analyzed** before,”  
thus could be a “good **assistive tool** for discovering new algorithms.”